



Reasoning under uncertainty


Ronojoy Adhikari
The Institute of Mathematical Sciences
Chennai, India

PyCon 2015
Bangalore, India







Logic
Probability



Logic
Probability



Probabilistic
programming



Logic
Probability

Probabilistic
programming

Probabilistic
networks

Logic
Probability



live code
scipy.stats

Probabilistic
programming

Probabilistic
networks

Logic
Probability



live code
scipy.stats

Probabilistic
programming



live code
lea

Probabilistic
networks

Logic
Probability



live code
scipy.stats

Probabilistic
programming



live code
lea

Probabilistic
networks



live code
lea
PyPNL ?

The purpose of computing is insight, not numbers.




The purpose of computing is insight, not numbers.



Richard Hamming

The purpose of computing is insight, not numbers.




What is the purpose of data science ?




What is the purpose of data science ?


Insight, not numbers!

Data science






Domain knowledge



Domain knowledge




Data curation



Domain knowledge



THE Data SCIENTIFIC METHOD




Data curation

Mathematical model

Domain knowledge



THE Data SCIENTIFIC METHOD



Data curation

Mathematical model


A/B testing

Domain knowledge



THE Data SCIENTIFIC METHOD

Machine learning



Data curation


Mathematical model

A/B testing

Domain knowledge


Machine
inference

Machine
learning



Domain knowledge

Value from data



1. Problem or Question



Data curation

7. Communicate Your Findings



THE Data SCIENTIFIC METHOD

Machine inference



6. Make a Conclusion

I was wrong!

I was right!

Try Again.
Make a New Hypothesis



Machine learning

Mathematical model



3. Make a Hypothesis



4. Experiment to Test Your Hypothesis

A/B testing

5. Gather and Analyze Data






© Patricia Felix






© Patricia Felix











© Patricia Felix














Question





Data



Inference



© Patricia Felix



Automated Inference ~ Machine Learning

Uncertainty

- how much will the Nile flood ?
- when will this equipment fail ?
- is this email spam ?
- is this applicant a good hire ?

Decisions

- should we invest in dams ?
- should we build redundancy ?
- should i delete without reading?
- should we look more ?

**We need to make reasoned decisions
in the face of uncertainty**

Reasoning : Logic - Boolean algebra

Uncertainty : Chance, probability

Combine : Bayesian probability

Logic and probability

Logic and probability

- Reasoning as the basis for a science of data
- Reasoning under certainty and under uncertainty
- Boolean logic and probability
- Rules of probability theory
- Assigning probabilities - indifference and maximum entropy
- Inference and learning
- Is this a fair coin ? Elementary example of reasoning under uncertainty

The scientific method

The scientific method




The scientific method



"Now this is the peculiarity of scientific method, that when once it has become a habit of mind, that mind converts all facts whatsoever into science. The field of science is unlimited; its solid contents are endless, every group of natural phenomena, every phase of social life, every stage of past or present development is material for science. The unity of all science consists alone in its method, not in its material.

The man who classifies facts of any kind whatever, who sees their mutual relation and describes their sequence, is applying the scientific method and is a man of science. The facts may belong to the past history of mankind, to the social statistics of our great cities, to the atmosphere of the most distant stars, to the digestive organs of a worm, or to the life of a scarcely visible bacillus. It is not the facts themselves which form science, but the method in which they are dealt with."


The scientific method



"Now this is the peculiarity of scientific method, that when once it has become a habit of mind, that mind converts all facts whatsoever into science. The field of science is unlimited; its solid contents are endless, every group of natural phenomena, every phase of social life, every stage of past or present development is material for science. The unity of all science consists alone in its method, not in its material.

The man who classifies facts of any kind whatever, who sees their mutual relation and describes their sequence, is applying the scientific method and is a man of science. The facts may belong to the past history of mankind, to the social statistics of our great cities, to the atmosphere of the most distant stars, to the digestive organs of a worm, or to the life of a scarcely visible bacillus. It is not the facts themselves which form science, but the method in which they are dealt with."

The scientific method




"Now this is the peculiarity of scientific method, that when once it has become a habit of mind, that mind converts all facts whatsoever into science. The field of science is unlimited; its solid contents are endless, every group of natural phenomena, every phase of social life, every stage of past or present development is material for science. The unity of all science consists alone in its method, not in its material.

The man who classifies facts of any kind whatever, who sees their mutual relation and describes their sequence, is applying the scientific method and is a man of science. The facts may belong to the past history of mankind, to the social statistics of our great cities, to the atmosphere of the most distant stars, to the digestive organs of a worm, or to the life of a scarcely visible bacillus. It is not the facts themselves which form science, but the method in which they are dealt with."


Reasoning in the scientific method

Deductive logic



Boolean algebra

Inductive logic



Bayesian probability

Boolean algebra : reasoning under certainty

- Formalization of Aristotelian logic
- Propositions : are either TRUE or FALSE
- Operations : conjunction (AND), disjunction (OR), negation (NOT)
- Laws : algebraic identities between compound propositions
- Ex.1 : $\text{NOT}(A \text{ AND } B) = (\text{NOT } A) \text{ OR } (\text{NOT } B)$
- Ex. 2 : $\text{NOT}(A \text{ OR } B) = (\text{NOT } A) \text{ AND } (\text{NOT } B)$
- Rules for reasoning consistently with certain propositions.

Probability : reasoning under uncertainty

- Generalization of Boolean logic
- Propositions have a truth value p , with $p = 0$ (FALSE) and $p = 1$ (TRUE)
- Operations : conjunction (AND), disjunction (OR), negation (NOT)
- sum rule : $P(A) + P(\text{NOT } A) = 1$
- product rule : $P(A \text{ AND } B) = P(A|B)P(B) = P(B|A) P(A)$
- $\Rightarrow P(A \text{ OR } B) = P(A) + P(B) - P(A \text{ AND } B)$
- independent $\Rightarrow P(A|B) = P(A)$; mutually exclusive $\Rightarrow P(A \text{ OR } B) = P(A) + P(B)$


Assigning probabilities

- Probabilities are **ALWAYS** conditioned on information $P(A) = P(A | I)$
- Consider a set of propositions A_1, A_2, \dots, A_n , that are exhaustive and mutually exclusive. In the absence of any other information, the **principle of indifference** says that $P(A_i) = 1/N$ (Laplace)
- When additional information is available, probabilities are assigned taking the additional information into account. The **principle of maximum entropy** says that P should be assigned by maximizing $\sum P_i \log P_i$, subject to the constraints that derive from the additional information.
- Maximum entropy reduces to indifference when there are no constraints.

Bayes theorem

- $P(A \text{ and } B) = P(A|B) P(B) = P(B|A)P(A)$
- $P(A|B) = P(B|A) P(A) / P(B)$
- Looks trivial but is extremely deep!
- $P(\text{disease} | \text{symptom})$ = want to know this.
- $P(\text{symptom} | \text{disease})$ = can estimate this (even empirically!)
- $P(\text{disease} | \text{symptom}) = P(\text{symptom} | \text{disease}) P(\text{disease}) / P(\text{symptom})$

Probabilistic networks


$$P(\text{disease} \mid \text{symptoms}) = \text{diagnosis}$$


RILACS

Representation

Inference

Learning

Actions



Bayesian inference : is this a fair coin ?

Bayesian inference : is this a fair coin ?

A coin is thrown 6 times with the outcome $D = \{hhhhtt\}$

Bayesian inference : is this a fair coin ?

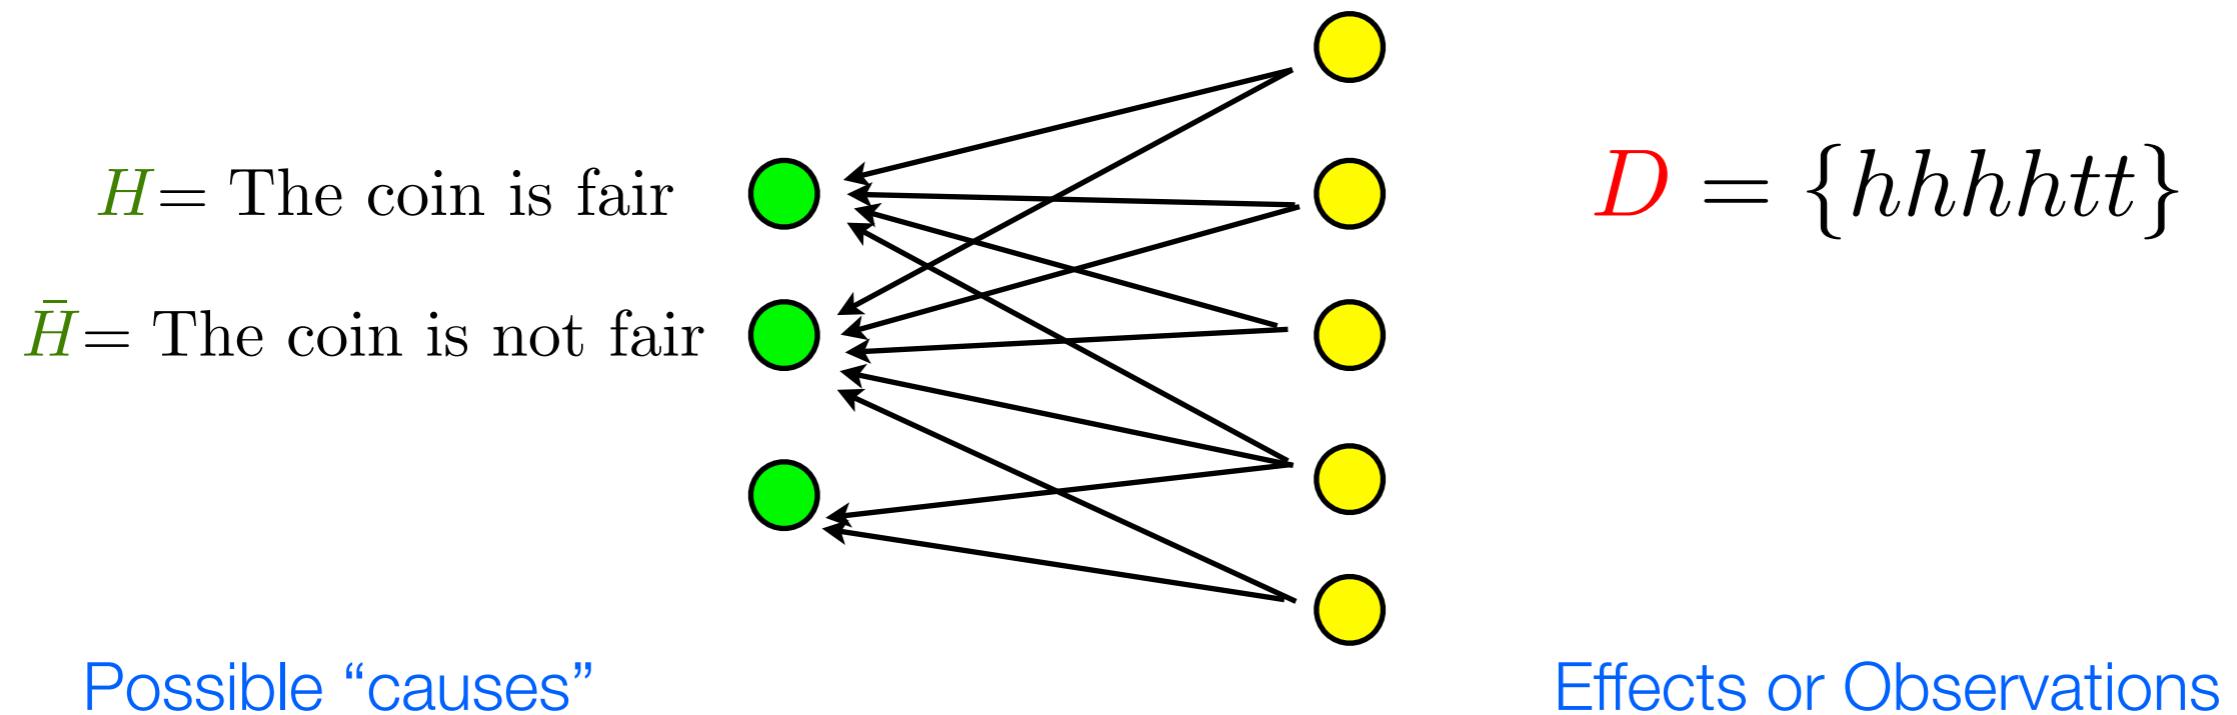
A coin is thrown 6 times with the outcome $D = \{hhhhtt\}$

What is the probability of the proposition $H =$ The coin is fair

Bayesian inference : is this a fair coin ?

A coin is thrown 6 times with the outcome $D = \{hhhhtt\}$

What is the probability of the proposition $H =$ The coin is fair



Bernoulli trials

$P(\textcolor{red}{D}|\textcolor{green}{H})$ = Probability of D *given* H
= Probability of outcome *given* a fair coin

$$P(\textcolor{red}{h}|\textcolor{green}{H}) = P(\textcolor{red}{t}|\textcolor{green}{H}) = \frac{1}{2}$$

$$P(\textcolor{red}{D}|\textcolor{green}{H}) = P(\textcolor{red}{hhhhtt}|\textcolor{green}{H}) = \frac{1}{2} \cdot \frac{1}{2} \cdot \frac{1}{2} \cdot \frac{1}{2} \cdot \frac{1}{2} \cdot \frac{1}{2}$$

$$P(\textcolor{red}{h}|\bar{H}) = \theta, P(\textcolor{red}{t}|\bar{H}) = 1 - \theta$$

$$P(\textcolor{red}{D}|\bar{H}) = P(\textcolor{red}{hhhhtt}|\bar{H}) = \theta \cdot \theta \cdot \theta \cdot \theta \cdot (1 - \theta) \cdot (1 - \theta)$$

Bernoulli trials

$$P(\textcolor{red}{D}|\theta) = P(\textcolor{red}{hhhhtt}|\theta) = \theta \cdot \theta \cdot \theta \cdot \theta \cdot (1 - \theta) \cdot (1 - \theta)$$

$$P(\theta|\textcolor{red}{D}) = P(\theta|\textcolor{red}{hhhhtt})$$

= Probability of the hypothesis θ given the data

$$\textcolor{green}{H} \rightarrow \theta = \frac{1}{2}$$

$$\bar{\textcolor{green}{H}} \rightarrow \theta \neq \frac{1}{2}$$

$$P(\textcolor{red}{D}|\theta)$$

know this

$$P(\theta|\textcolor{red}{D})$$

want this

Bernoulli trials

$$P(\theta | \textcolor{red}{D}) = \frac{P(\textcolor{red}{D} | \theta) P(\theta)}{P(\textcolor{red}{D})}$$

Bernoulli trials

$$P(n_1|\theta, N) = \frac{N!}{n_1!(N-n_1)!} \theta^{n_1} (1-\theta)^{N-n_1}$$

P(D|H) - likelihood


$$P(\theta) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} \theta^{a-1} (1-\theta)^{b-1}$$

P(H) = prior $\langle \theta \rangle = \frac{a}{a+b}$

$$P(\theta|n_1, N) \sim \theta^{n_1+a-1} (1-\theta)^{n-n_1+b-1}$$

P(H|D) = posterior


Bernoulli trials



Posterior distribution with uniform prior


Live coding 1

Probabilistic programming



Ronald Fisher

Let the data speak for themselves!



Ronald Fisher

Let the data speak for themselves!



Edwin Jaynes

The data *cannot* speak for themselves;
and they never have, in any real problem
of inference.

Classification | Clustering

Regression

Dimensionality reduction

Classification

predict class, given attributes

Clustering


Regression

Dimensionality reduction

Classification

predict class, given attributes

Animal Characteristics
Copyright 1998 Norsys Software Corp.




Clustering

Regression Dimensionality reduction

Classification

predict class, given attributes

Animal Characteristics
Copyright 1998 Norsys Software Corp.



Clustering

Regression


predict values, given other values

Dimensionality reduction

Classification

predict class, given attributes


Animal Characteristics
Copyright 1998 Norsys Software Corp.



Clustering

Regression

predict values, given other values



"It's a non-linear pattern with outliers....but for some reason I'm very happy with the data."


Dimensionality reduction

Classification

predict class, given attributes

Animal Characteristics

Copyright 1998 Norsys Software Corp.



Clustering

group similar things together

Regression

predict values, given other values

Dimensionality reduction




"It's a non-linear pattern with outliers....but for some reason I'm very happy with the data."

Classification

predict class, given attributes


Animal Characteristics

Copyright 1998 Norsys Software Corp.




Clustering

group similar things together



Regression

predict values, given other values



"It's a non-linear pattern with outliers....but for some reason I'm very happy with the data."


Dimensionality reduction

Classification

predict class, given attributes


Animal Characteristics

Copyright 1998 Norsys Software Corp.




Clustering

group similar things together



Regression

predict values, given other values



"It's a non-linear pattern with outliers....but for some reason I'm very happy with the data."

Dimensionality reduction


keeping only the relevant variables

Classification

predict class, given attributes


Animal Characteristics

Copyright 1998 Norsys Software Corp.




Clustering

group similar things together



Regression


predict values, given other values



"It's a non-linear pattern with outliers....but for some reason I'm very happy with the data."

Dimensionality reduction

keeping only the relevant variables



Hypotheses
require
mathematical models

Bayesian | Frequentist

Blackbox

Causal

Bayesian

probability is a state of knowledge

Frequentist

Blackbox

Causal

Bayesian

probability is a state of knowledge


Frequentist

probability is a frequency

Blackbox

Causal

Bayesian
probability is a state of knowledge




Frequentist
probability is a frequency

Blackbox

Causal

Bayesian
probability is a state of knowledge




Frequentist
probability is a frequency

Blackbox
ML : toolbox for processing data


Causal

Bayesian
probability is a state of knowledge

Frequentist
probability is a frequency




Blackbox
ML : toolbox for processing data



Bayesian

probability is a state of knowledge




Frequentist

probability is a frequency

Blackbox

ML : toolbox for processing data




Causal

ML : learning generative models of data


Bayesian

probability is a state of knowledge



Blackbox

ML : toolbox for processing data




Frequentist

probability is a frequency

Causal

ML : learning generative models of data




Bayesian

probability is a state of knowledge



Causal

ML : learning generative models of data



Bayesian


probability is a state of knowledge



Nearly all approaches to probabilistic programming are Bayesian since it is hard to create other coherent frameworks for automated reasoning about uncertainty - Ghahramani, Nature, 2015

Causal

ML : learning generative models of data





Machine Learning is Ubiquitous and Very Useful

ISR



DARPA Grand Challenge

Fully autonomous ground vehicles competition

Nuclear Test Ban Treaty Compliance: Deduce set of seismic events given detections and misdetections on a network of stations

Image Search/Activity Detection: Find and identify objects and actions in video

Object Tracking: Follow vehicles as they move through a city and are recorded in multiple video streams (DARPA CZTS)

Patterns of Life: Process wide area aerial surveillance data and associated tracks to infer location types and object dynamics

Bird Migration Patterns: Model spatio-temporal distribution of birds (by species); involves large-scale sensor integration

DARPA LAGR: Vision-based robot navigation

Google Glasses: Perform searches based on images taken by user cell phone cameras

Natural Language Processing



Siri

Voice recognition and Natural Language Processing (NLP)

Watson: Computer system capable of answering questions posed in natural language

Topic Models: Statistical model for discovering the abstract "topics" that occur in a collection of documents

Distributed Topic Models: Asynchronous distributed topic discovery

Citation Analysis: Given citations, extract author, title, and venue strings and identify co-reference citations

Entity Resolution: Discovering entities that are mentioned in a large corpus of news articles

NLP Sequence Tagging: Tagging parts of speech and recognizing named entities in newspaper articles

Predictive Analytics



Netflix Challenge
Predict user ratings for films based on previous ratings

Microsoft Matchbox: Match players based on their gaming skill set

Predictive Database: Understand information based on causal relationships in data

Bing Image Search: Search for images on the web by selecting text in word document

Amazon Recommendation Engine: Recommend items based on consumer data

Cyber and Other



ORNL's Attack Variant Detector: Discover compromised systems


Yahoo's Bayesian Spam Filtering: Self-adapting system based on word probabilities

Cyber Genome Lineage: Reverse engineer malware samples to find shared "genetic" features between different malware samples


Gene Sequencing: Determine order of nucleotides in a DNA molecule

Disclaimer: Images of specific products are used for illustration only. Use of these images does not imply endorsement of inherent technical vulnerabilities.

- Brittleness of implementations & lack of reusable tools.
*PHY decoder: \$100M/standard;
Infer.NET: max 20% on model*
- High level of required expertise
*10K solvers,
100s of grad student hours per model*
- Painfully slow & unpredictable solvers
Massive data sets, complex algorithms, tricky coding for graph traversal and numeric stability
- Challenges constructing models
*Limited modeling vocabulary;
models entwined with solvers*



We're Missing a Tool to Write These Applications



The Missing Tool (Explained by Example)

Model :

```

mem strength person = gaussian 100 10
lazy person = flip 0.1
pulling person = if lazy person then (strength person) / 2
                           else strength person
total-pulling team = sum (map pulling team)
winner team1 team2 = greater (total-pulling team1)
                           (total-pulling team2)

```



Source: Noah Goodman, POPL (2013)

Query :

strength Bob


System will calculate probability distribution
for Bob's strength given known facts

Facts :

[Bob, Mark] = winner [Bob, Mark] [Tom, Sam]
 [Bob, Fred] = winner [Bob, Fred] [Jon, Jim]

The user describes the model at a high level. An inference engine analyzes the program, query, data, and available hardware resources to produce best solution


What is happening in the Python space ?



Live coding 2

Probabilistic networks


Probabilistic network : conditional independences



causal reasoning : $P(\text{tired} \mid \text{flu})$

diagnostic reasoning : $P(\text{flu} \mid \text{tired})$

Markov blanket



Live coding 3

This repository Search

Pull requests Issues Gist

ValueFromData / PyPNL

Unwatch 4

Python Probabilistic Network Library — Edit

4 commits 1 branch 0 releases 1 contributor

Branch: master + PyPNL / +

add shorter copy

ronojoy authored 2 minutes ago latest commit 689586f250

README.md add shorter copy 2 minutes ago


README.md

PyPNL

PyPNL is a probabilistic network library in Python. The library consists of a Python API which interfaces to state of the art C/C++ inference and learning engines. The API design is modular and allows for plug-ins which augment functionality. PyPNL comes with documentation, examples, and a test suite.

Features

The PyPNL library has a rich set of features for **representation, inference, learning and actions** using probabilistic networks.



This repository Search

Pull requests Issues Gist

ValueFromData / PyPNL

Unwatch 4

Python Probabilistic Network Library — Edit

4 commits 1 branch 0 releases 1 contributor

Branch: master PyPNL / +

add shorter copy ronojoy authored 2 minutes ago latest commit 689586f250 2 minutes ago

README.md add shorter copy


README.md

PyPNL

PyPNL is a probabilistic network library in Python. The library consists of a Python API which interfaces to state of the art C/C++ inference and learning engines. The API design is modular and allows for plug-ins which augment functionality. PyPNL comes with documentation, examples, and a test suite.

Features

The PyPNL library has a rich set of features for **representation, inference, learning and actions** using probabilistic networks.



We are building a causal learning and inference engine that will beat the current state-of-art!

This repository Search

Pull requests Issues Gist

ValueFromData / PyPNL

Unwatch 4

Python Probabilistic Network Library — Edit

4 commits 1 branch 0 releases 1 contributor

Branch: master PyPNL / +

add shorter copy ronojoy authored 2 minutes ago latest commit 689586f250 2 minutes ago

README.md add shorter copy


README.md

PyPNL

PyPNL is a probabilistic network library in Python. The library consists of a Python API which interfaces to state of the art C/C++ inference and learning engines. The API design is modular and allows for plug-ins which augment functionality. PyPNL comes with documentation, examples, and a test suite.

Features

The PyPNL library has a rich set of features for **representation, inference, learning and actions** using probabilistic networks.



We are building a causal learning and inference engine that will beat the current state-of-art!

This repository Search

Pull requests Issues Gist

ValueFromData / PyPNL

Unwatch 4

Python Probabilistic Network Library — Edit

4 commits 1 branch 0 releases 1 contributor

Branch: master PyPNL / +

add shorter copy ronojoy authored 2 minutes ago latest commit 689586f250 2 minutes ago

README.md add shorter copy

README.md

PyPNL

PyPNL is a probabilistic network library in Python. The library consists of a Python API which interfaces to state of the art C/C++ inference and learning engines. The API design is modular and allows for plug-ins which augment functionality. PyPNL comes with documentation, examples, and a test suite.

Features

The PyPNL library has a rich set of features for **representation, inference, learning and actions** using probabilistic networks.

Thank you for your attention!